

PolicyWatcher: Evidence-preserving monitoring of online terms and privacy policies

Fabrizio Degni^{a,*}

^aPOLIMI Graduate School of Management, Via Lambruschini 4/c, Building 26/a, 20156 Milan, Italy

Abstract

PolicyWatcher is open-source software for continuously retrieving, validating, versioning, and communicating changes to online terms of service and privacy policies. Its central design goal is to prevent plausible but unsupported data from reaching public outputs. A five-path acquisition cascade combines direct HTTP, explicit HTTP/2, optional browser rendering, and two timestamp-guarded web archives. Retrieved pages pass content validation and provenance checks before normalized text, hashes, diffs, and structured bilingual analyses are stored. A shared public-evidence gate excludes development fixtures, partial retrievals, and unverified records from pages, APIs, reports, search indexes, and subscriber notifications. In a dated operational snapshot of 50 configured documents from 16 providers, 46 records passed the public-evidence gate and four were suspended; a fixture-only control exposed no policy content or change records. The software supports reproducible research on policy accessibility, retrieval-path dependence, longitudinal change, and evidence-aware civic technology while explicitly separating operational observations from claims about legal or semantic accuracy.

Keywords: policy monitoring, web measurement, digital governance, provenance, civic technology, research software

Nr.	Code metadata description	Please fill in this column
C1	Current code version	v3.6.4, SoftwareX snapshot 1171b06
C2	Permanent GitHub link to code/repository used for this code version	GitHub commit 1171b06
C3	Legal Code License	MIT License
C4	Code versioning system used	Git
C5	Software code languages, tools, and services used	TypeScript, JavaScript, Next.js, React, Prisma, SQLite, Node.js, Playwright/Chromium, Google Gemini API
C6	Compilation requirements, operating environments, and dependencies	Node.js 20 or later; npm; dependencies declared in <code>package-lock.json</code> ; Chromium is needed only for the optional remote renderer
C7	Link to developer documentation/manual	README at commit 1171b06
C8	Support email for questions	fabrizio.degni@gsom.polimi.it

Table 1: Code metadata (mandatory).

Required Metadata

Current code version

1. Motivation and significance

Terms of service and privacy policies define how online services collect, retain, transfer, and reuse information. Yet the documents are long, change over time, and are rarely read in full. Prior work has quantified the reading burden of privacy policies [1], documented user acceptance without meaningful reading [2], and created annotated or longitudinal corpora for privacy research [3, 4]. Systems such as Polisis and CLAUDETTE analyze policy language [5, 6], while ToS;DR and TOSBack provide public summaries or tracked versions [7, 8].

These contributions leave a practical research-software gap. Retrospective corpora can fix their source set in advance, but a live monitor repeatedly faces bot challenges, JavaScript-only pages, regional variants, URL changes, consent interstitials, and archives that may be older than the current baseline.

*Corresponding author

Email address: `fabrizio.degni@gsom.polimi.it` (Fabrizio Degni)

URL: `https://www.policywatcher.online` (Fabrizio Degni)

The dominant failure is often not a crash. It is a plausible wrong document, such as a challenge page or stale archive, that is stored and analyzed as if it were the provider’s policy.

POLICYWATCHER was developed to make that evidence boundary explicit. It monitors a configured set of public legal documents and records how each observation was obtained. It provides a running web application for citizens and practitioners, but it is also research software: its retrieval logs, normalized versions, provenance states, and public-evidence decisions can support studies of policy accessibility, longitudinal change, and measurement reliability. The software does not present model-generated explanations as legal advice, and its operational measurements are not treated as a benchmark of legal interpretation accuracy.

The user can deploy the application on commodity hosting, configure provider URLs and jurisdictions, schedule bounded scans, inspect source-quality findings, and publish only records that meet the evidence contract. Readers access a bilingual public timeline, per-document diffs, comparison views, evidence-status notices, and downloadable reports. Administrators additionally receive check logs, dataset-quality findings, review history, and explicit controls for source promotion or suspension.

2. Software description

2.1. Software architecture

POLICYWATCHER is a server-rendered TypeScript application built with Next.js and React. Prisma provides typed access to SQLite, while a small authenticated rendering service can execute Chromium on infrastructure where a full browser runtime is available. The hosted language model is an optional downstream analysis component; acquisition, validation, hashing, evidence status, and public gating do not depend on a successful model response.

Figure 1 shows the evidence path. A scheduler submits each configured URL to the acquisition cascade. The result is validated before any text can enter the evidence store. Normalized text feeds hashing and diffing so cosmetic markup changes do not become policy changes. Public outputs query a shared gate rather than duplicating visibility logic at each page or API.

2.2. Acquisition and content contract

The acquisition cascade orders five strategies by fidelity to the live source: direct HTTP/1.1, explicit HTTP/2, optional headless rendering, the Wayback Machine [9], and Common Crawl [10]. Direct requests use bounded retries,

PolicyWatcher evidence-preserving architecture

Every public record must retain retrieval provenance




Figure 1: PolicyWatcher architecture. Retrieval provenance and evidence status travel with each record; uncertain records follow the withholding path.

body-aware timeouts, manual redirect handling, and address validation at every hop. HTTP/2 is a separate path for servers that reject non-browser HTTP/1.1 clients. The renderer executes script-dependent pages but remains optional. Archives are witnesses of last resort.

Archive fallback is governed by a freshness rule: a candidate capture is admissible only when its archive timestamp is not earlier than the last live baseline. A stale capture therefore cannot replace a newer live document and create a false regression. Captures without a usable timestamp are rejected.

Every retrieval returns exactly one of three states: `ok`, `invalid`, or `unavailable`. An `ok` result includes validated text, final URL, retrieval path, hash, and timing metadata. `invalid` means that bytes arrived but failed content checks; `unavailable` means that no admissible path returned evidence. Callers must handle all states explicitly.

Content validation detects challenge pages, soft not-found responses, maintenance and login pages, implausibly short bodies, and redirects that drift to another registrable domain or to a generic site root. Boilerplate and hidden elements are removed before deterministic normalization. Because the fetcher operates server-side, URLs and every redirect also pass server-side request forgery controls based on scheme restrictions, address classification, DNS validation, and connection-time address pinning [11].

2.3. Evidence storage, analysis, and governance

For valid text, the system compares a SHA-256 hash with the current baseline. An unchanged observation creates a check-log record. A changed observation creates an immutable snapshot and a complete line diff. The optional analysis layer then requests bilingual summaries, key points, risk rationales, regional impact, and a fifteen-indicator governance assessment. Model output is parsed against a strict schema; out-of-range values or unrecognized categories cause rejection rather than silent repair. Retrieved policy text is treated as untrusted model context, and generated output can never substitute for source text.

Development fixtures are represented separately from evidence. Each policy has an ingestion method and lifecycle state; snapshots and changes carry an explicit public-evidence flag. The public gate admits only available or reviewed, non-seeded records with evidence. The same predicate protects pages, APIs, reports, sitemap entries, share images, and email notifications. Suspended records expose only a minimal source-status notice.

A controlled re-baseline operation promotes a fixture to its first retrieved baseline without emitting a fabricated change. Promotion is allowed only when the record is seeded, has no verified check-log evidence, and has no snapshot already marked public. Administrative decisions and ignored quality findings are appended to an audit log with author and rationale.

2.4. Software functionalities

The main functions are: scheduled and scoped source checking; evidence-aware version storage; full-text diffing; bilingual structured analysis; source suspension and remediation; a fifteen-indicator comparison matrix; per-company and per-policy evidence views; public change feeds and reports; subscriber notifications; dataset-assurance checks; source onboarding; encrypted export; and renderer health monitoring. Continuous integration runs schema checks, build and lint tasks, automated tests, dependency auditing, static analysis, and the dataset-assurance command. Version 3.6.4 contains 145 automated tests across 24 test modules, including pure tests for validation, freshness, evidence-gate, and destructive-operation predicates.

3. Illustrative examples

Figure 2 follows one scheduled document check. Suppose a configured privacy-policy URL has a stored live baseline. A valid direct or rendered response is normalized and hashed. If the hash is unchanged, the system

records the successful check without publishing a change. If it differs, the system stores the new snapshot and diff, requests analysis, and publishes the change only after the evidence gate succeeds.

If the response is a consent wall, challenge page, short placeholder, or redirected home page, the result is **invalid**; the page text never becomes a baseline. If live methods fail, an archive may be used only when its capture time is at least as recent as the previous live baseline. If all paths fail, the result is **unavailable**. Public pages may then display a source verification notice but cannot display policy text, a risk score, a change event, or a notification generated from the failed check.

Illustrative policy-check workflow

One scheduled document check; three honest outcomes




Figure 2: One document check produces evidence, a diagnosed invalid response, or explicit unavailability. Only the evidence branch can reach publication.

A fixture-only installation provides a second example. Its configured records have realistic shapes for interface development but are marked as seeded and lack retrieved evidence. Querying the same public APIs used in production yields zero public policies, zero changes, and suspension notices for every configured record. This control exercises the boundary independently of provider network conditions.

4. Impact

The principal impact is methodological: POLICYWATCHER makes acquisition uncertainty a queryable part of the dataset instead of an exception hidden in scanner logs. Researchers can separate configured coverage from

evidence coverage, compare retrieval-path dependence across providers, examine the duration and causes of source unavailability, and construct longitudinal corpora whose records retain source, capture time, normalization path, and publication decision. The same design supports experiments on change-detection accuracy and human agreement without conflating a failed fetch with a negative observation.

Table 2 reports a dated, reproducible operational snapshot. On 11 July 2026, the running deployment contained 50 configured documents from 16 providers. Forty-six records passed the public-evidence gate and four were suspended. The public inventory contained 21 change records. Among the 46 public records, 40 used direct retrieval, five used the renderer, and one used an archive path. The fixture control exposed no policy or change content. The JSON snapshot and its read-only collection script are supplied with the submission.

Operational measure	Production	Fixture control
Configured policy records	50	50
Records admitted as public evidence	46 (92%)	0 (0%)
Suspended records	4 (8%)	50 (100%)
Publicly listed providers	15 of 16	0 of 16
Public change-record inventory	21	0

Table 2: Public operational state captured on 11 July 2026. Inventory counts are not estimates of semantic or legal accuracy.

The operational deployment shows that the fallback paths are exercised rather than merely specified. It also demonstrates that withholding is observable: unavailable documents remain counted and diagnosable instead of disappearing from the denominator. The software runs on shared application hosting plus a small renderer service, so longitudinal collection is feasible without a large institutional cluster.

For research practice, POLICYWATCHER can contribute in four ways. First, it can produce timestamped provider-version corpora for policy and digital-governance studies. Second, its per-attempt logs support controlled ablations of direct, HTTP/2, rendered, and archive retrieval. Third, its explicit negative states allow measurement of source accessibility and censorship-by-friction without treating blocked pages as missing at random. Fourth, its normalized diffs and evidence links can support blinded human annotation of substantive changes and evaluation of downstream summaries.

The current evidence does not establish precision or recall for semantic change detection, agreement with legal experts, or correctness of generated

risk annotations. There is not yet a validated claim of broad external adoption, and the 21 public change records are inventory items rather than independently adjudicated legal events. These boundaries define the next evaluations: a pre-registered repeated-scan window, a versioned ground-truth corpus, cascade ablation, at least two independent annotators, and reporting of agreement, precision, recall, false-change rate, and time to detection.

The project is publicly developed on GitHub [12], provides installation and contribution documentation, and holds a passing OpenSSF Best Practices self-assessment [13]. These are engineering-quality and reusability signals, not security or legal certifications. The software is designed for extension through new providers, policy types, jurisdictions, retrieval strategies, public views, and analysis backends.

5. Conclusions

POLICYWATCHER provides an open implementation of evidence-aware monitoring for online terms and privacy policies. Its reusable contribution is the combination of a multi-path acquisition cascade, deterministic content contract, freshness rule, provenance-preserving version store, and shared public-evidence gate. Together, these mechanisms keep retrieval failures and development fixtures from becoming apparently authoritative public records. The software is already deployed and supplies reproducible operational evidence, while its explicit limitations leave legal and semantic validity to properly designed expert and longitudinal studies.

Data availability

The operational snapshot supporting Table 2, together with the read-only collection script and interpretation notes, is supplied as Supplementary Data and in the submission source archive. The live public endpoints are available at <https://www.policywatcher.online>. No private check logs, subscriber information, credentials, administrative data, or raw provider text are included.

CRedit author statement

Fabrizio Degni: Conceptualization, Methodology, Software, Validation, Investigation, Data curation, Visualization, Writing - original draft, Writing - review and editing, Project administration.

Funding

This research received no specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of competing interest

Fabrizio Degni is the creator and maintainer of PolicyWatcher and operates the public demonstration service. The author declares no financial interest that could have appeared to influence the work reported in this paper.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work, the author used OpenAI Codex to assist with editorial restructuring, language refinement, LaTeX preparation, and generation of reproducible vector-figure source. The author reviewed, verified, and edited the resulting material and takes full responsibility for the content of the publication.

Acknowledgements

The author thanks the open-source maintainers and public-interest archival services on which the software depends.

References

- [1] A. M. McDonald, L. F. Cranor, The cost of reading privacy policies, *I/S: A Journal of Law and Policy for the Information Society* 4 (3) (2008) 543-568.
- [2] J. A. Obar, A. Oeldorf-Hirsch, The biggest lie on the internet: ignoring the privacy policies and terms of service policies of social networking services, *Information, Communication & Society* 23 (1) (2020) 128-147. doi:10.1080/1369118X.2018.1486870.
- [3] S. Wilson, F. Schaub, A. A. Dara, F. Liu, S. Cherivirala, P. G. Leon, et al., The creation and analysis of a website privacy policy corpus, in: *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*, 2016, pp. 1330-1340. doi:10.18653/v1/P16-1126.

- [4] R. Amos, G. Acar, E. Lucherini, M. Kshirsagar, A. Narayanan, J. Mayer, Privacy policies over time: curation and analysis of a million-document dataset, in: Proceedings of The Web Conference, 2021, pp. 2165-2176. doi:10.1145/3442381.3450048.
- [5] H. Harkous, K. Fawaz, R. Lebrete, F. Schaub, K. G. Shin, K. Aberer, Polisis: automated analysis and presentation of privacy policies using deep learning, in: Proceedings of the 27th USENIX Security Symposium, 2018, pp. 531-548.
- [6] M. Lippi, P. Palka, G. Contissa, F. Lagioia, H.-W. Micklitz, G. Sartor, P. Torroni, CLAUDETTE: an automated detector of potentially unfair clauses in online terms of service, Artificial Intelligence and Law 27 (2) (2019) 117-139. doi:10.1007/s10506-019-09243-2.
- [7] Terms of Service; Didn't Read, ToS;DR public service and data, <https://tosdr.org>, accessed 23 July 2026.
- [8] Electronic Frontier Foundation, TOSBack: the terms-of-service tracker, <https://tosback.org>, accessed 23 July 2026.
- [9] Internet Archive, Wayback Machine, <https://web.archive.org>, accessed 23 July 2026.
- [10] Common Crawl Foundation, Common Crawl corpus, <https://commoncrawl.org>, accessed 23 July 2026.
- [11] Open Worldwide Application Security Project, Server Side Request Forgery Prevention Cheat Sheet, https://cheatsheetseries.owasp.org/cheatsheets/Server_Side_Request_Forgery_Prevention_Cheat_Sheet.html, accessed 23 July 2026.
- [12] F. Degni, PolicyWatcher version 3.6.4 [computer software], GitHub, 2026, GitHub commit 1171b06.
- [13] OpenSSF Best Practices Badge Program, PolicyWatcher project 13465, <https://www.bestpractices.dev/projects/13465>, accessed 23 July 2026.

Current executable software version

The public deployment is available at <https://www.policywatcher.online>. Installation and execution instructions for the reproducible source release are provided in the repository linked in Table 1.